

Stepwise Linear Regression

Veerasak Kritsanapraphan

15/1/2020

Stepwise Linear Regression

Import packages necessary first.

```
library(MASS)
library(plyr)
library(ggplot2)
library(knitr)
```

Prepare data

```
# Assign better variable names
colnames(birthwt) <- c("birthwt.below.2500", "mother.age", "mother.weight",
  "race", "mother.smokes", "previous.prem.labor", "hypertension", "uterine.irr",
  "physician.visits", "birthwt.grams")

# Assign better labels to categorical variables
birthwt <- transform(birthwt,
  race = as.factor(mapvalues(race, c(1, 2, 3),
    c("white", "black", "other"))),
  mother.smokes = as.factor(mapvalues(mother.smokes,
    c(0,1), c("no", "yes"))),
  hypertension = as.factor(mapvalues(hypertension,
    c(0,1), c("no", "yes"))),
  uterine.irr = as.factor(mapvalues(uterine.irr,
    c(0,1), c("no", "yes"))),
  birthwt.below.2500 = as.factor(mapvalues(birthwt.below.2500,
    c(0,1), c("no", "yes")))
)
```

Run logistic regression

```
formula = birthwt.grams ~ mother.age + mother.weight + physician.visits + mother.smokes + uterine.irr +
fullmod = lm(formula, data = birthwt)
summary(fullmod)
```

```
##
## Call:
## lm(formula = formula, data = birthwt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1689.4  -467.6    -7.6    464.5   1750.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2508.467    294.477   8.518 6.15e-15 ***
## mother.age         4.745      9.721   0.488 0.626019
## mother.weight     4.272      1.725   2.477 0.014172 *
```

```
## physician.visits      -8.012      48.113  -0.167  0.867923
## mother.smokesyes     -228.486     102.506  -2.229  0.027044 *
## uterine.irryes       -527.097     143.887  -3.663  0.000327 ***
## previous.prem.labor   -71.709     105.515  -0.680  0.497619
## hypertensionyes      -642.048     209.323  -3.067  0.002492 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 674.2 on 181 degrees of freedom
## Multiple R-squared:  0.1769, Adjusted R-squared:  0.1451
## F-statistic: 5.559 on 7 and 181 DF,  p-value: 8.177e-06
```

No independent variables

```
nothing <- lm(birthwt.grams ~ 1,family=binomial, data = birthwt)
```

```
## Warning: In lm.fit(x, y, offset = offset, singular.ok = singular.ok, ...) :
## extra argument 'family' will be disregarded
```

```
summary(nothing)
```

```
##
## Call:
## lm(formula = birthwt.grams ~ 1, data = birthwt, family = binomial)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -2235.59  -530.59   32.41   542.41  2045.41
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2944.59      53.04   55.51  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 729.2 on 188 degrees of freedom
```

Stepwise Linear Regression (“Backward”)

```
backwards = step(fullmod) # Backwards selection is the default
```

```
## Start:  AIC=2469.96
## birthwt.grams ~ mother.age + mother.weight + physician.visits +
##      mother.smokes + uterine.irr + previous.prem.labor + hypertension
##
##              Df Sum of Sq      RSS      AIC
## - physician.visits    1      12607 82293521 2468.0
## - mother.age          1      108335 82389248 2468.2
## - previous.prem.labor  1       209962 82490875 2468.4
## <none>                        82280914 2470.0
## - mother.smokes       1      2258613 84539526 2473.1
## - mother.weight       1      2788878 85069791 2474.3
## - hypertension       1      4276845 86557758 2477.5
## - uterine.irr         1      6100389 88381303 2481.5
##
## Step:  AIC=2467.99
## birthwt.grams ~ mother.age + mother.weight + mother.smokes +
```

```

##      uterine.irr + previous.prem.labor + hypertension
##
##              Df Sum of Sq      RSS      AIC
## - mother.age      1      98286 82391807 2466.2
## - previous.prem.labor 1      206657 82500178 2466.5
## <none>                                82293521 2468.0
## - mother.smokes      1      2257167 84550688 2471.1
## - mother.weight      1      2783045 85076566 2472.3
## - hypertension      1      4273665 86567186 2475.6
## - uterine.irr      1      6089392 88382912 2479.5
##
## Step:  AIC=2466.21
## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +
##      previous.prem.labor + hypertension
##
##              Df Sum of Sq      RSS      AIC
## - previous.prem.labor 1      175821 82567628 2464.6
## <none>                                82391807 2466.2
## - mother.smokes      1      2314920 84706727 2469.4
## - mother.weight      1      3103917 85495724 2471.2
## - hypertension      1      4383189 86774996 2474.0
## - uterine.irr      1      6244871 88636678 2478.0
##
## Step:  AIC=2464.62
## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +
##      hypertension
##
##              Df Sum of Sq      RSS      AIC
## <none>                                82567628 2464.6
## - mother.smokes      1      2623742 85191369 2468.5
## - mother.weight      1      3306358 85873986 2470.0
## - hypertension      1      4441221 87008848 2472.5
## - uterine.irr      1      6979875 89547503 2477.9

```

```
formula(backwards)
```

```

## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +
##      hypertension

```

```
summary(backwards)
```

```

##
## Call:
## lm(formula = birthwt.grams ~ mother.weight + mother.smokes +
##      uterine.irr + hypertension, data = birthwt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1665.02  -452.62    11.16   473.64  1858.65
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2577.096     226.794   11.363 < 2e-16 ***
## mother.weight         4.506       1.660    2.714 0.007270 **
## mother.smokesyes   -242.113     100.127   -2.418 0.016579 *

```

```
## uterine.irryes    -549.878    139.424   -3.944 0.000114 ***
## hypertensionyes  -649.098    206.327   -3.146 0.001931 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 669.9 on 184 degrees of freedom
## Multiple R-squared:  0.1741, Adjusted R-squared:  0.1561
## F-statistic: 9.695 on 4 and 184 DF,  p-value: 3.886e-07
```

Stepwise Linear Regression (“Forward”)

```
forwards <- step(fullmod, direction="forward")
```

```
## Start:  AIC=2469.96
## birthwt.grams ~ mother.age + mother.weight + physician.visits +
##      mother.smokes + uterine.irr + previous.prem.labor + hypertension
```

```
formula(forwards)
```

```
## birthwt.grams ~ mother.age + mother.weight + physician.visits +
##      mother.smokes + uterine.irr + previous.prem.labor + hypertension
```

```
summary(forwards)
```

```
##
## Call:
## lm(formula = birthwt.grams ~ mother.age + mother.weight + physician.visits +
##      mother.smokes + uterine.irr + previous.prem.labor + hypertension,
##      data = birthwt)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1689.4  -467.6    -7.6    464.5   1750.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    2508.467    294.477   8.518 6.15e-15 ***
## mother.age         4.745      9.721   0.488 0.626019
## mother.weight      4.272      1.725   2.477 0.014172 *
## physician.visits   -8.012     48.113  -0.167 0.867923
## mother.smokesyes  -228.486    102.506  -2.229 0.027044 *
## uterine.irryes    -527.097    143.887  -3.663 0.000327 ***
## previous.prem.labor -71.709    105.515  -0.680 0.497619
## hypertensionyes   -642.048    209.323  -3.067 0.002492 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 674.2 on 181 degrees of freedom
## Multiple R-squared:  0.1769, Adjusted R-squared:  0.1451
## F-statistic: 5.559 on 7 and 181 DF,  p-value: 8.177e-06
```

Stepwise Linear Regression (“Both”)

```
both <- step(fullmod, direction="both")
```

```
## Start:  AIC=2469.96
## birthwt.grams ~ mother.age + mother.weight + physician.visits +
##      mother.smokes + uterine.irr + previous.prem.labor + hypertension
```

```

##
##              Df Sum of Sq      RSS      AIC
## - physician.visits      1      12607 82293521 2468.0
## - mother.age            1      108335 82389248 2468.2
## - previous.prem.labor    1       209962 82490875 2468.4
## <none>                    82280914 2470.0
## - mother.smokes         1      2258613 84539526 2473.1
## - mother.weight         1      2788878 85069791 2474.3
## - hypertension          1      4276845 86557758 2477.5
## - uterine.irr           1      6100389 88381303 2481.5
##
## Step:  AIC=2467.99
## birthwt.grams ~ mother.age + mother.weight + mother.smokes +
##      uterine.irr + previous.prem.labor + hypertension
##
##              Df Sum of Sq      RSS      AIC
## - mother.age            1       98286 82391807 2466.2
## - previous.prem.labor    1      206657 82500178 2466.5
## <none>                    82293521 2468.0
## + physician.visits      1      12607 82280914 2470.0
## - mother.smokes         1      2257167 84550688 2471.1
## - mother.weight         1      2783045 85076566 2472.3
## - hypertension          1      4273665 86567186 2475.6
## - uterine.irr           1      6089392 88382912 2479.5
##
## Step:  AIC=2466.21
## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +
##      previous.prem.labor + hypertension
##
##              Df Sum of Sq      RSS      AIC
## - previous.prem.labor    1      175821 82567628 2464.6
## <none>                    82391807 2466.2
## + mother.age            1       98286 82293521 2468.0
## + physician.visits      1       2559 82389248 2468.2
## - mother.smokes         1      2314920 84706727 2469.4
## - mother.weight         1      3103917 85495724 2471.2
## - hypertension          1      4383189 86774996 2474.0
## - uterine.irr           1      6244871 88636678 2478.0
##
## Step:  AIC=2464.62
## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +
##      hypertension
##
##              Df Sum of Sq      RSS      AIC
## <none>                    82567628 2464.6
## + previous.prem.labor    1      175821 82391807 2466.2
## + mother.age            1       67449 82500178 2466.5
## + physician.visits      1       2121 82565507 2466.6
## - mother.smokes         1      2623742 85191369 2468.5
## - mother.weight         1      3306358 85873986 2470.0
## - hypertension          1      4441221 87008848 2472.5
## - uterine.irr           1      6979875 89547503 2477.9

```

```
formula(both)
```

```
## birthwt.grams ~ mother.weight + mother.smokes + uterine.irr +  
## hypertension
```

```
summary(both)
```

```
##  
## Call:  
## lm(formula = birthwt.grams ~ mother.weight + mother.smokes +  
## uterine.irr + hypertension, data = birthwt)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max   
## -1665.02  -452.62    11.16   473.64  1858.65   
##  
## Coefficients:  
##              Estimate Std. Error t value Pr(>|t|)      
## (Intercept)   2577.096    226.794   11.363 < 2e-16 ***  
## mother.weight     4.506      1.660    2.714 0.007270 **  
## mother.smokesyes -242.113    100.127  -2.418 0.016579 *  
## uterine.irryes   -549.878    139.424  -3.944 0.000114 ***  
## hypertensionyes -649.098    206.327  -3.146 0.001931 **  
## ---  
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1  
##  
## Residual standard error: 669.9 on 184 degrees of freedom  
## Multiple R-squared:  0.1741, Adjusted R-squared:  0.1561  
## F-statistic: 9.695 on 4 and 184 DF, p-value: 3.886e-07
```